

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import VotingClassifier
```

```
# Load the training dataset
train_data = pd.read_csv('/content/Churn_TRAIN.csv')
```

```
# Display the first few rows of the training dataset
print(train_data.head())
```

```
# Load the test dataset
test_data = pd.read_csv('/content/Churn_TEST.csv')
```

```
# Display the first few rows of the test dataset
print(test_data.head())
```

	Call	Failure	Complains	Subscription	Length	Charge	Amount	\
0		10	0		37		1	
1		4	0		36		0	
2		2	0		9		1	
3		14	0		22		5	
4		6	0		25		0	

	Seconds of Use	Frequency of use	Frequency of SMS	\
0	6908	76	117	
1	15295	182	175	
2	4390	40	215	
3	3238	53	48	
4	395	19	3	

	Distinct Called Numbers	Age Group	Tariff Plan	Status	Age	\
0	40	2	1	1	25	
1	32	2	1	1	25	
2	10	3	1	1	30	
3	25	3	2	1	30	
4	5	2	1	2	25	

	Customer Value	Churn
0	840.780	0
1	1483.965	0
2	1037.200	0
3	323.640	0
4	32.130	0

	Call	Failure	Complains	Subscription	Length	Charge	Amount	\
0		4	0		42		0	
1		3	0		10		2	
2		13	0		27		0	
3		5	0		42		0	
4		19	0		19		2	

	Seconds of Use	Frequency of use	Frequency of SMS	\
0	2315	43	293	
1	2593	35	16	
2	945	28	12	
3	888	17	24	
4	6453	144	90	

	Distinct Called Numbers	Age Group	Tariff Plan	Status	Age	\
0	22	3	1	1	30	
1	13	3	1	1	30	
2	7	3	1	2	30	
3	10	4	1	2	45	
4	44	2	2	1	25	

	Customer Value	Churn
0	1266.320	0
1	169.120	0
2	86.920	0
3	82.625	1
4	701.865	0

```
# Separate features and target variable in the training dataset
X_train = train_data.drop(columns=['Churn'])
y_train = train_data['Churn']
```

```
# Separate features and target variable in the test dataset
```

```
X_test = test_data.drop(columns=['Churn'])
y_test = test_data['Churn']
```

```
# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
# Train a RandomForestClassifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
```

```
# Train a LogisticRegression
lr = LogisticRegression(max_iter=1000)
lr.fit(X_train, y_train)
```

```
# Train a SVC
svc = SVC(probability=True)
svc.fit(X_train, y_train)
```

```

      SVC
SVC(probability=True)

```

```
# Make predictions using the RandomForestClassifier
rf_predictions = rf.predict(X_test)
```

```
# Evaluate the RandomForestClassifier
print('Classification Report for RandomForestClassifier:')
print(classification_report(y_test, rf_predictions))
print('Confusion Matrix for RandomForestClassifier:')
print(confusion_matrix(y_test, rf_predictions))
```

```
# Make predictions using the LogisticRegression
lr_predictions = lr.predict(X_test)
```

```
# Evaluate the LogisticRegression
print('Classification Report for LogisticRegression:')
print(classification_report(y_test, lr_predictions))
print('Confusion Matrix for LogisticRegression:')
print(confusion_matrix(y_test, lr_predictions))
```

```
# Make predictions using the SVC
svc_predictions = svc.predict(X_test)
```

```
# Evaluate the SVC
print('Classification Report for SVC:')
print(classification_report(y_test, svc_predictions))
print('Confusion Matrix for SVC:')
print(confusion_matrix(y_test, svc_predictions))
```

```
Classification Report for RandomForestClassifier:
              precision    recall  f1-score   support

     0       0.96       0.98       0.97       1324
     1       0.89       0.77       0.83        251

   accuracy       0.95       0.95       0.95       1575
  macro avg       0.92       0.88       0.90       1575
weighted avg       0.95       0.95       0.95       1575
```

```
Confusion Matrix for RandomForestClassifier:
[[1299  25]
 [ 57 194]]
```

```
Classification Report for LogisticRegression:
              precision    recall  f1-score   support

     0       0.90       0.97       0.94       1324
     1       0.76       0.43       0.55        251

   accuracy       0.89       0.89       0.89       1575
  macro avg       0.83       0.70       0.74       1575
weighted avg       0.88       0.89       0.87       1575
```

```
Confusion Matrix for LogisticRegression:
[[1290  34]
 [ 143 108]]
```

```
Classification Report for SVC:
              precision    recall  f1-score   support
```

0	0.92	0.99	0.95	1324
1	0.94	0.54	0.68	251
accuracy			0.92	1575
macro avg	0.93	0.77	0.82	1575
weighted avg	0.92	0.92	0.91	1575

Confusion Matrix for SVC:

```
[[1315  9]
 [ 116 135]]
```

```
# Hyperparameter tuning for RandomForestClassifier
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='roc_auc')
grid_search.fit(X_train, y_train)
print('Best parameters for RandomForestClassifier:', grid_search.best_params_)
print('Best score for RandomForestClassifier:', grid_search.best_score_)
```

```
# Ensemble learning using VotingClassifier
voting_clf = VotingClassifier(estimators=[
    ('lr', LogisticRegression(max_iter=1000)),
    ('rf', RandomForestClassifier()),
    ('svc', SVC(probability=True))],
    voting='soft')
voting_clf.fit(X_train, y_train)
voting_predictions = voting_clf.predict(X_test)
print('Classification Report for VotingClassifier:')
print(classification_report(y_test, voting_predictions))
print('Confusion Matrix for VotingClassifier:')
print(confusion_matrix(y_test, voting_predictions))
```

```
# Handling imbalanced data using SMOTE
smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)
```

```
Best parameters for RandomForestClassifier: {'max_depth': 30, 'min_samples_split': 5, 'n_estimators': 300}
Best score for RandomForestClassifier: 0.980417347446427
Classification Report for VotingClassifier:
```

	precision	recall	f1-score	support
0	0.93	0.99	0.96	1324
1	0.90	0.62	0.73	251
accuracy			0.93	1575
macro avg	0.91	0.80	0.84	1575
weighted avg	0.93	0.93	0.92	1575

Confusion Matrix for VotingClassifier:

```
[[1306  18]
 [  96 155]]
```

```
# Import necessary libraries
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_auc_score, roc_curve, precision_recall_curve, auc
```

```
# Calculate the ROC AUC scores
rf_roc_auc = roc_auc_score(y_test, rf.predict(X_test))
lr_roc_auc = roc_auc_score(y_test, lr.predict(X_test))
svc_roc_auc = roc_auc_score(y_test, svc.predict(X_test))
voting_roc_auc = roc_auc_score(y_test, voting_clf.predict(X_test))
```

```
print('ROC AUC score for RandomForestClassifier: ', rf_roc_auc)
print('ROC AUC score for LogisticRegression: ', lr_roc_auc)
print('ROC AUC score for SVC: ', svc_roc_auc)
print('ROC AUC score for VotingClassifier: ', voting_roc_auc)
```

```
ROC AUC score for RandomForestClassifier: 0.8770130956536392
ROC AUC score for LogisticRegression: 0.7022995630769973
ROC AUC score for SVC: 0.7655255112480591
ROC AUC score for VotingClassifier: 0.8019673571574727
```


```
# Plot the ROC curves
plt.figure(figsize=(7, 7))
```

```
for model, name, color in zip([rf, lr, svc, voting_clf], ['RandomForest', 'Logistic Regression', 'SVC', 'Voting'], ['blue', 'green', 'red', 'purple']):
    fpr, tpr, _ = roc_curve(y_test, model.predict_proba(X_test)[:, 1])
    plt.plot(fpr, tpr, color=color, label=name)
```

```

plt.plot([0, 1], [0, 1], color='navy', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

```



```

# Cost analysis
# Please replace the values for cost_fp and cost_fn with your actual costs
cost_fp = 1 # Cost of a false positive
cost_fn = 1 # Cost of a false negative

cm_rf = confusion_matrix(y_test, rf_predictions)
cm_lr = confusion_matrix(y_test, lr_predictions)
cm_svc = confusion_matrix(y_test, svc_predictions)
cm_voting = confusion_matrix(y_test, voting_predictions)

cost_rf = cm_rf[0, 1] * cost_fp + cm_rf[1, 0] * cost_fn
cost_lr = cm_lr[0, 1] * cost_fp + cm_lr[1, 0] * cost_fn
cost_svc = cm_svc[0, 1] * cost_fp + cm_svc[1, 0] * cost_fn
cost_voting = cm_voting[0, 1] * cost_fp + cm_voting[1, 0] * cost_fn

print('Total cost for RandomForestClassifier: ', cost_rf)
print('Total cost for LogisticRegression: ', cost_lr)
print('Total cost for SVC: ', cost_svc)
print('Total cost for VotingClassifier: ', cost_voting)

Total cost for RandomForestClassifier: 82
Total cost for LogisticRegression: 177
Total cost for SVC: 125
Total cost for VotingClassifier: 114

# Import necessary libraries
from sklearn.metrics import precision_score, recall_score, f1_score, log_loss, matthews_corrcoef

# Calculate metrics for RandomForestClassifier
rf_precision = precision_score(y_test, rf_predictions)
rf_recall = recall_score(y_test, rf_predictions)
rf_f1 = f1_score(y_test, rf_predictions)
rf_log_loss = log_loss(y_test, rf.predict_proba(X_test))
rf_matthews = matthews_corrcoef(y_test, rf_predictions)

print('RandomForestClassifier:')

```

```

print('Precision: ', rf_precision)
print('Recall: ', rf_recall)
print('F1-score: ', rf_f1)
print('Log Loss: ', rf_log_loss)
print('Matthews Correlation Coefficient: ', rf_matthews)

# Calculate metrics for LogisticRegression
lr_precision = precision_score(y_test, lr_predictions)
lr_recall = recall_score(y_test, lr_predictions)
lr_f1 = f1_score(y_test, lr_predictions)
lr_log_loss = log_loss(y_test, lr.predict_proba(X_test))
lr_matthews = matthews_corrcoef(y_test, lr_predictions)

print('\nLogisticRegression:')
print('Precision: ', lr_precision)
print('Recall: ', lr_recall)
print('F1-score: ', lr_f1)
print('Log Loss: ', lr_log_loss)
print('Matthews Correlation Coefficient: ', lr_matthews)

# Calculate metrics for SVC
svc_precision = precision_score(y_test, svc_predictions)
svc_recall = recall_score(y_test, svc_predictions)
svc_f1 = f1_score(y_test, svc_predictions)
svc_log_loss = log_loss(y_test, svc.predict_proba(X_test))
svc_matthews = matthews_corrcoef(y_test, svc_predictions)

print('\nSVC:')
print('Precision: ', svc_precision)
print('Recall: ', svc_recall)
print('F1-score: ', svc_f1)
print('Log Loss: ', svc_log_loss)
print('Matthews Correlation Coefficient: ', svc_matthews)

# Calculate metrics for VotingClassifier
voting_precision = precision_score(y_test, voting_predictions)
voting_recall = recall_score(y_test, voting_predictions)
voting_f1 = f1_score(y_test, voting_predictions)
voting_log_loss = log_loss(y_test, voting_clf.predict_proba(X_test))
voting_matthews = matthews_corrcoef(y_test, voting_predictions)

print('\nVotingClassifier:')
print('Precision: ', voting_precision)
print('Recall: ', voting_recall)
print('F1-score: ', voting_f1)
print('Log Loss: ', voting_log_loss)
print('Matthews Correlation Coefficient: ', voting_matthews)

RandomForestClassifier:
Precision: 0.8858447488584474
Recall: 0.7729083665338645
F1-score: 0.825531914893617
Log Loss: 0.220508289503508
Matthews Correlation Coefficient: 0.7976555339516986

LogisticRegression:
Precision: 0.7605633802816901
Recall: 0.4302788844621514
F1-score: 0.549618320610687
Log Loss: 0.2300518912407888
Matthews Correlation Coefficient: 0.5170571241823865

SVC:
Precision: 0.9375
Recall: 0.5378486055776892
F1-score: 0.6835443037974683
Log Loss: 0.21601221478887933
Matthews Correlation Coefficient: 0.6743976367376933

VotingClassifier:
Precision: 0.8959537572254336
Recall: 0.6175298804780877
F1-score: 0.7311320754716981
Log Loss: 0.1665534415242046
Matthews Correlation Coefficient: 0.7069257293672487

# Convert numpy arrays back into dataframes
X_train = pd.DataFrame(X_train, columns=train_data.drop(columns=['Churn']).columns)
X_test = pd.DataFrame(X_test, columns=train_data.drop(columns=['Churn']).columns)

```

```

from sklearn.inspection import permutation_importance
import matplotlib.pyplot as plt

# Calculate permutation importance for RandomForestClassifier
result = permutation_importance(rf, X_test, y_test, n_repeats=10, random_state=0, n_jobs=-1)

# Sort features by importance
sorted_idx = result.importances_mean.argsort()

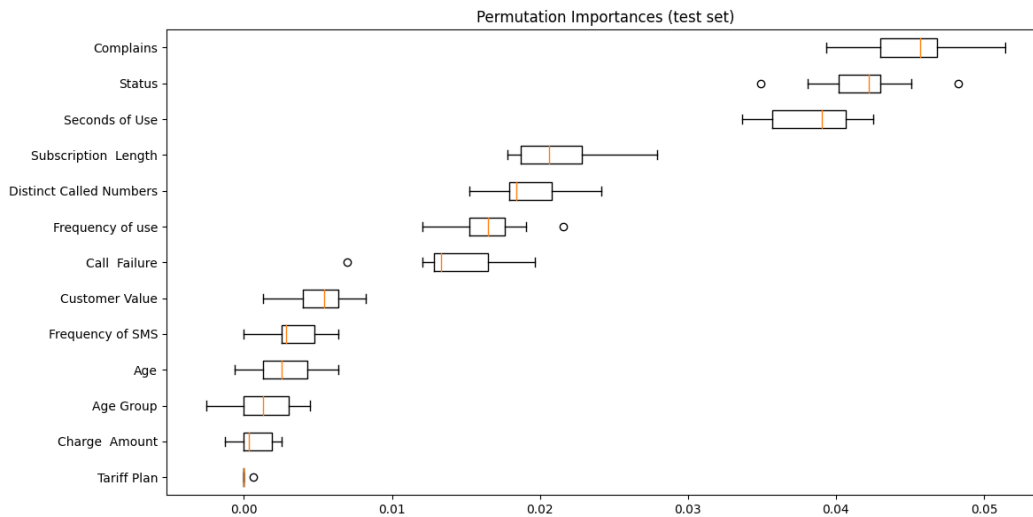
# Plot
fig, ax = plt.subplots(figsize=(12, 6))
ax.boxplot(result.importances[sorted_idx].T, vert=False, labels=X_test.columns[sorted_idx])
ax.set_title("Permutation Importances (test set)")
fig.tight_layout()
plt.show()

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but RandomForestClassifier does not
warnings.warn(

```



```

# Define cost values
cost_tp = -100 # Cost of true positive (retention cost)
cost_fp = -100 # Cost of false positive (unnecessary retention cost)
cost_tn = 0 # Cost of true negative (no cost)
cost_fn = -2300 # Cost of false negative (lost revenue + cost of acquiring a new customer)

# Calculate total cost for RandomForestClassifier
cost_rf = cm_rf[0, 1] * cost_fp + cm_rf[1, 0] * cost_fn + cm_rf[1, 1] * cost_tp + cm_rf[0, 0] * cost_tn

# Calculate total cost for LogisticRegression
cost_lr = cm_lr[0, 1] * cost_fp + cm_lr[1, 0] * cost_fn + cm_lr[1, 1] * cost_tp + cm_lr[0, 0] * cost_tn

# Calculate total cost for SVC
cost_svc = cm_svc[0, 1] * cost_fp + cm_svc[1, 0] * cost_fn + cm_svc[1, 1] * cost_tp + cm_svc[0, 0] * cost_tn

# Calculate total cost for VotingClassifier
cost_voting = cm_voting[0, 1] * cost_fp + cm_voting[1, 0] * cost_fn + cm_voting[1, 1] * cost_tp + cm_voting[0, 0] * cost_tn

print('Total cost for RandomForestClassifier: ', cost_rf)
print('Total cost for LogisticRegression: ', cost_lr)
print('Total cost for SVC: ', cost_svc)
print('Total cost for VotingClassifier: ', cost_voting)

```

```

Total cost for RandomForestClassifier: -153000
Total cost for LogisticRegression: -343100
Total cost for SVC: -281200
Total cost for VotingClassifier: -238100

```

```

# Define profit values
profit_tn = 1000 # Profit from true negative (customer correctly identified as not churning)
profit_tp = 1000 # Profit from true positive (customer correctly identified as churning and retained)

# Calculate total profit for RandomForestClassifier
profit_rf = cm_rf[0, 0] * profit_tn + cm_rf[1, 1] * profit_tp

# Calculate total profit for LogisticRegression
profit_lr = cm_lr[0, 0] * profit_tn + cm_lr[1, 1] * profit_tp

# Calculate total profit for SVC
profit_svc = cm_svc[0, 0] * profit_tn + cm_svc[1, 1] * profit_tp

# Calculate total profit for VotingClassifier
profit_voting = cm_voting[0, 0] * profit_tn + cm_voting[1, 1] * profit_tp

print('Total profit for RandomForestClassifier: ', profit_rf)
print('Total profit for LogisticRegression: ', profit_lr)
print('Total profit for SVC: ', profit_svc)
print('Total profit for VotingClassifier: ', profit_voting)

Total profit for RandomForestClassifier: 1493000
Total profit for LogisticRegression: 1398000
Total profit for SVC: 1450000
Total profit for VotingClassifier: 1461000

import numpy as np

from sklearn.model_selection import cross_val_score

# Perform cross-validation on RandomForestClassifier
rf_cv_scores = cross_val_score(rf, X_train, y_train, cv=5)
print('Cross-validation scores for RandomForestClassifier: ', rf_cv_scores)
print('Mean cross-validation score for RandomForestClassifier: ', np.mean(rf_cv_scores))

# Perform cross-validation on LogisticRegression
lr_cv_scores = cross_val_score(lr, X_train, y_train, cv=5)
print('\nCross-validation scores for LogisticRegression: ', lr_cv_scores)
print('Mean cross-validation score for LogisticRegression: ', np.mean(lr_cv_scores))

# Perform cross-validation on SVC
svc_cv_scores = cross_val_score(svc, X_train, y_train, cv=5)
print('\nCross-validation scores for SVC: ', svc_cv_scores)
print('Mean cross-validation score for SVC: ', np.mean(svc_cv_scores))

# Perform cross-validation on VotingClassifier
voting_cv_scores = cross_val_score(voting_clf, X_train, y_train, cv=5)
print('\nCross-validation scores for VotingClassifier: ', voting_cv_scores)
print('Mean cross-validation score for VotingClassifier: ', np.mean(voting_cv_scores))

Cross-validation scores for RandomForestClassifier: [0.95555556 0.93650794 0.95238095 0.93333333 0.95555556]
Mean cross-validation score for RandomForestClassifier: 0.9466666666666667

Cross-validation scores for LogisticRegression: [0.9047619 0.9015873 0.87619048 0.85714286 0.92063492]
Mean cross-validation score for LogisticRegression: 0.8920634920634921

Cross-validation scores for SVC: [0.92063492 0.93015873 0.89206349 0.88253968 0.93333333]
Mean cross-validation score for SVC: 0.9117460317460317

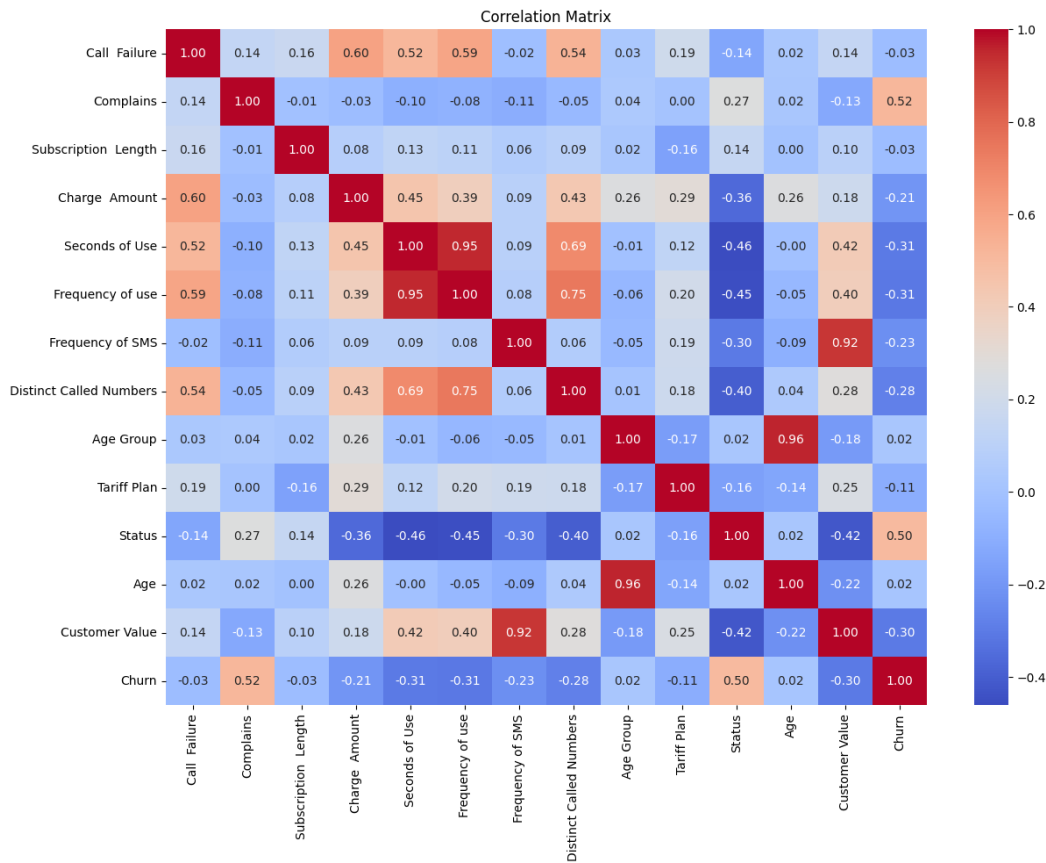
Cross-validation scores for VotingClassifier: [0.94920635 0.93015873 0.92698413 0.8984127 0.94285714]
Mean cross-validation score for VotingClassifier: 0.9295238095238094

import seaborn as sns

# Calculate correlation matrix
corr = train_data.corr()

# Plot the heatmap
plt.figure(figsize=(14, 10))
sns.heatmap(corr, annot=True, fmt=".2f", cmap='coolwarm', cbar=True)
plt.title('Correlation Matrix')
plt.show()

```



```

from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
# Define the model with n_init specified
kmeans = KMeans(n_clusters=3, n_init=10, random_state=0)

# Define data for clustering (only numerical variables)
cluster_data = train_data.select_dtypes(include=[np.number])

# Scale the data
scaler = StandardScaler()
cluster_data_scaled = scaler.fit_transform(cluster_data)

# Define the model
kmeans = KMeans(n_clusters=3, random_state=0) # we choose 3 clusters, but this could be any number

# Fit the model
kmeans.fit(cluster_data_scaled)

# Get cluster labels
train_data['cluster'] = kmeans.labels_

# Check the size of each cluster
print(train_data['cluster'].value_counts())

# Explore the mean values in each cluster
cluster_summary = train_data.groupby('cluster').mean()
print(cluster_summary)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change fr
warnings.warn(
0      878
2      363
1      334

```



```
Name: cluster, dtype: int64
```

	Call	Failure	Complains	Subscription	Length	Charge	Amount	\
cluster								
0		5.446469	0.119590		32.082005		0.381549	
1		15.209581	0.029940		33.979042		2.407186	
2		6.093664	0.019284		31.016529		0.947658	

	Seconds of Use	Frequency of use	Frequency of SMS	\
cluster				
0	2089.438497	38.136674	13.514806	
1	10603.736527	151.011976	50.538922	
2	5003.107438	75.451791	241.440771	

	Distinct Called Numbers	Age Group	Tariff Plan	Status	Age	\
cluster						
0	15.480638	2.769932	1.001139	1.439636	30.461276	
1	44.224551	3.245509	1.095808	1.000000	35.284431	
2	23.988981	2.490358	1.253444	1.000000	27.410468	

	Customer Value	Churn
cluster		
0	137.852215	0.275626
1	575.416317	0.000000
2	1222.589518	0.005510

```
from sklearn.model_selection import learning_curve
```

```
# Define a function to plot the learning curves
```

```
def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None, n_jobs=None):
```

```
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()

    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
    return plt
```

```
# Generate learning curves
```

```
plot_learning_curve(rf, 'Learning curve for RandomForestClassifier', X_train, y_train, cv=5)
```

```
<module 'matplotlib.pyplot' from '/usr/local/lib/python3.10/dist-packages/matplotlib/pyplot.py'>
```

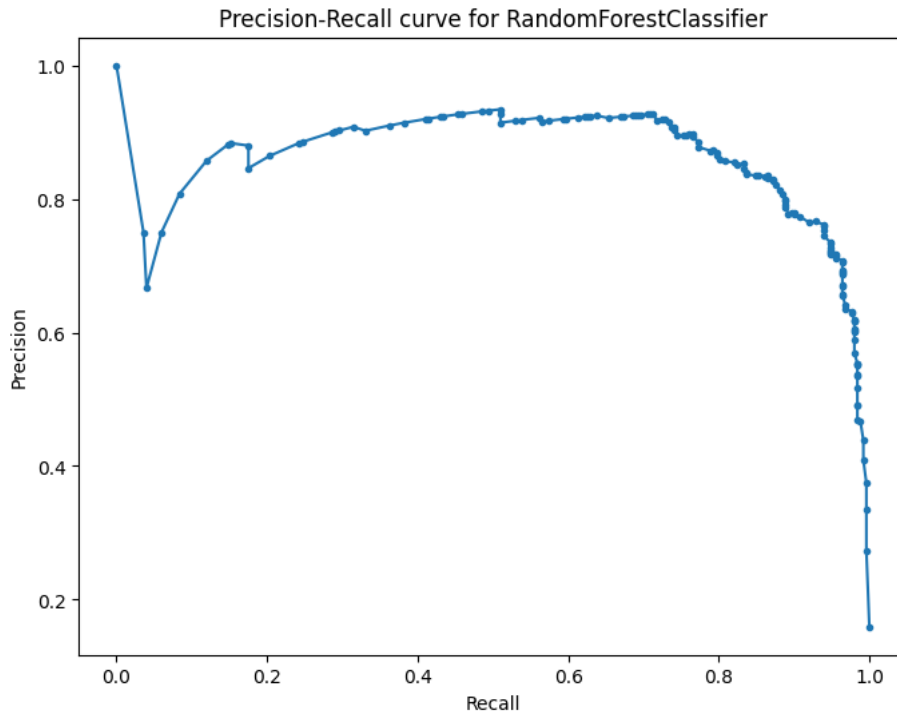
### Learning curve for RandomForestClassifier

```
from sklearn.metrics import precision_recall_curve

# Calculate precision and recall
precision, recall, _ = precision_recall_curve(y_test, rf.predict_proba(X_test)[: , 1])

# Plot the precision-recall curve
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall curve for RandomForestClassifier')
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but Ranc
warnings.warn(
```



```
# Binning the 'Age' column into three categories
train_data['Age_bin'] = pd.cut(train_data['Age'], bins=[0, 30, 60, 100], labels=['Young', 'Middle-aged', 'Senior'])
```

```
from imblearn.over_sampling import SMOTE

# Initialize SMOTE
smote = SMOTE()

# Fit SMOTE to the data
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

import xgboost as xgb

# Initialize XGBClassifier
xgb_clf = xgb.XGBClassifier()

# Train the model
xgb_clf.fit(X_train, y_train)

# Make predictions
y_pred_xgb = xgb_clf.predict(X_test)

# Evaluate the model
print(classification_report(y_test, y_pred_xgb))
```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	1324

1	0.88	0.78	0.83	251
accuracy			0.95	1575
macro avg	0.92	0.88	0.90	1575
weighted avg	0.95	0.95	0.95	1575

```
# Make predictions using multiple models
y_pred_rf = rf.predict_proba(X_test)[: , 1]
y_pred_lr = lr.predict_proba(X_test)[: , 1]
y_pred_xgb = xgb_clf.predict_proba(X_test)[: , 1]
```

```
# Average the predictions
y_pred_ensemble = (y_pred_rf + y_pred_lr + y_pred_xgb) / 3
```

```
# Convert probabilities to class labels
y_pred_ensemble = [1 if prob > 0.5 else 0 for prob in y_pred_ensemble]
```

```
# Evaluate the ensemble
print(classification_report(y_test, y_pred_ensemble))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but RandomForestClassifier was fitted without feature names
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has feature names, but LogisticRegression was fitted without feature names
warnings.warn(
              precision    recall  f1-score   support

         0       0.95      0.98      0.97       1324
         1       0.88      0.75      0.81        251

   accuracy          0.92
  macro avg          0.92
 weighted avg          0.94
```

```
import joblib
joblib.dump(rf, 'model.joblib')
```

```
['model.joblib']
```

```
!pip install shap
```

```
Collecting shap
  Downloading shap-0.42.0-cp310-cp310-manylinux_2_12_x86_64.manylinux2010_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (54.1 MB)
    547.1/547.1 kB 14.2 MB/s eta 0:00:00
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from shap) (1.22.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from shap) (1.10.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (from shap) (1.2.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from shap) (1.5.3)
Requirement already satisfied: tqdm>=4.27.0 in /usr/local/lib/python3.10/dist-packages (from shap) (4.65.0)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.10/dist-packages (from shap) (23.1)
Collecting slicer==0.0.7 (from shap)
  Downloading slicer-0.0.7-py3-none-any.whl (14 kB)
Requirement already satisfied: numba in /usr/local/lib/python3.10/dist-packages (from shap) (0.56.4)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.10/dist-packages (from shap) (2.2.1)
Requirement already satisfied: llvmlite<0.40,>=0.39.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba->shap) (0.39.1)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from numba->shap) (67.7.2)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->shap) (2022.7.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (1.3.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn->shap) (3.1.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1->pandas->shap) (1.16.0)
Installing collected packages: slicer, shap
Successfully installed shap-0.42.0 slicer-0.0.7
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
# Define your RandomForestClassifier
rf = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
# Fit the classifier to your training data
# (Substitute X_train and y_train with your actual training data)
rf.fit(Churn_TRAIN, Churn_TEST)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-7-5c1ce8ec3bbf> in <cell line: 8>()
      6 # Fit the classifier to your training data
      7 # (Substitute X_train and y_train with your actual training data)
----> 8 rf.fit(Churn_TRAIN, Churn_TEST)

NameError: name 'Churn_TRAIN' is not defined
```

Churn Train Overview