

```
!pip install scikit-learn
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, roc_auc_score
```

```
# Load the datasets
train_data = pd.read_csv('/content/Churn_TRAIN.csv')
test_data = pd.read_csv('/content/Churn_TEST.csv')
```

```
# Split the data into input features (X) and target variable (y)
X_train = train_data.drop('Churn', axis=1)
y_train = train_data['Churn']
X_test = test_data.drop('Churn', axis=1)
y_test = test_data['Churn']
```

```
# Train a random forest classifier
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
```

```
# Make predictions on the test set
y_pred_rf = rf.predict(X_test)
```

```
# Compute and print the classification report
print(classification_report(y_test, y_pred_rf))
```

```
# Compute and print the ROC AUC
roc_auc_rf = roc_auc_score(y_test, y_pred_rf)
print('ROC AUC (Random Forest):', roc_auc_rf)
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.22.4)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.10.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.1.0)
```


	precision	recall	f1-score	support
0	0.96	0.98	0.97	1324
1	0.89	0.78	0.83	251
accuracy			0.95	1575
macro avg	0.92	0.88	0.90	1575
weighted avg	0.95	0.95	0.95	1575

```
ROC AUC (Random Forest): 0.8829891912711691
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
plt.figure(figsize=(10, 6))
sns.histplot(data=train_data, x='Age', bins=30, kde=True)
plt.title('Age Distribution')
plt.show()
```


Age Distribution



```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, roc_auc_score
```

```
plt.figure(figsize=(10, 6))
sns.countplot(x='Churn', data=train_data)
plt.title('Churn Distribution')
plt.show()
```

Churn Distribution



```
X_train = train_data.drop('Churn', axis=1)
y_train = train_data['Churn']
X_test = test_data.drop('Churn', axis=1)
y_test = test_data['Churn']
```

```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

```
LogisticRegression
LogisticRegression(max_iter=1000)
```

```
y_pred = model.predict(X_test)
```

```
print(classification_report(y_test, y_pred))
print('ROC AUC:', roc_auc_score(y_test, y_pred))
```

```
precision    recall  f1-score   support

0           0.90      0.97      0.94       1324
```

	1	0.76	0.43	0.55	251
accuracy				0.89	1575
macro avg	0.83	0.70	0.74		1575
weighted avg	0.88	0.89	0.87		1575

ROC AUC: 0.7022995630769973

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report, roc_auc_score, confusion_matrix
from sklearn.model_selection import cross_val_score
```

```
# Preprocess your data
X_train = train_data.drop('Churn', axis=1)
y_train = train_data['Churn']
X_test = test_data.drop('Churn', axis=1)
y_test = test_data['Churn']

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train and evaluate different models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "Support Vector Machine": SVC()
}
```

```
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"Model: {name}")
    print(classification_report(y_test, y_pred))
    print('ROC AUC:', roc_auc_score(y_test, y_pred))
    print("-----")
```

Model: Logistic Regression					
	precision	recall	f1-score	support	
0	0.90	0.97	0.94	1324	
1	0.76	0.43	0.55	251	
accuracy			0.89	1575	
macro avg	0.83	0.70	0.74	1575	
weighted avg	0.88	0.89	0.87	1575	

ROC AUC: 0.7022995630769973

Model: Random Forest					
	precision	recall	f1-score	support	
0	0.96	0.98	0.97	1324	
1	0.88	0.79	0.83	251	
accuracy			0.95	1575	
macro avg	0.92	0.89	0.90	1575	
weighted avg	0.95	0.95	0.95	1575	

ROC AUC: 0.8858403245025938

Model: Support Vector Machine					
	precision	recall	f1-score	support	
0	0.92	0.99	0.95	1324	
1	0.94	0.54	0.68	251	
accuracy			0.92	1575	
macro avg	0.93	0.77	0.82	1575	
weighted avg	0.92	0.92	0.91	1575	

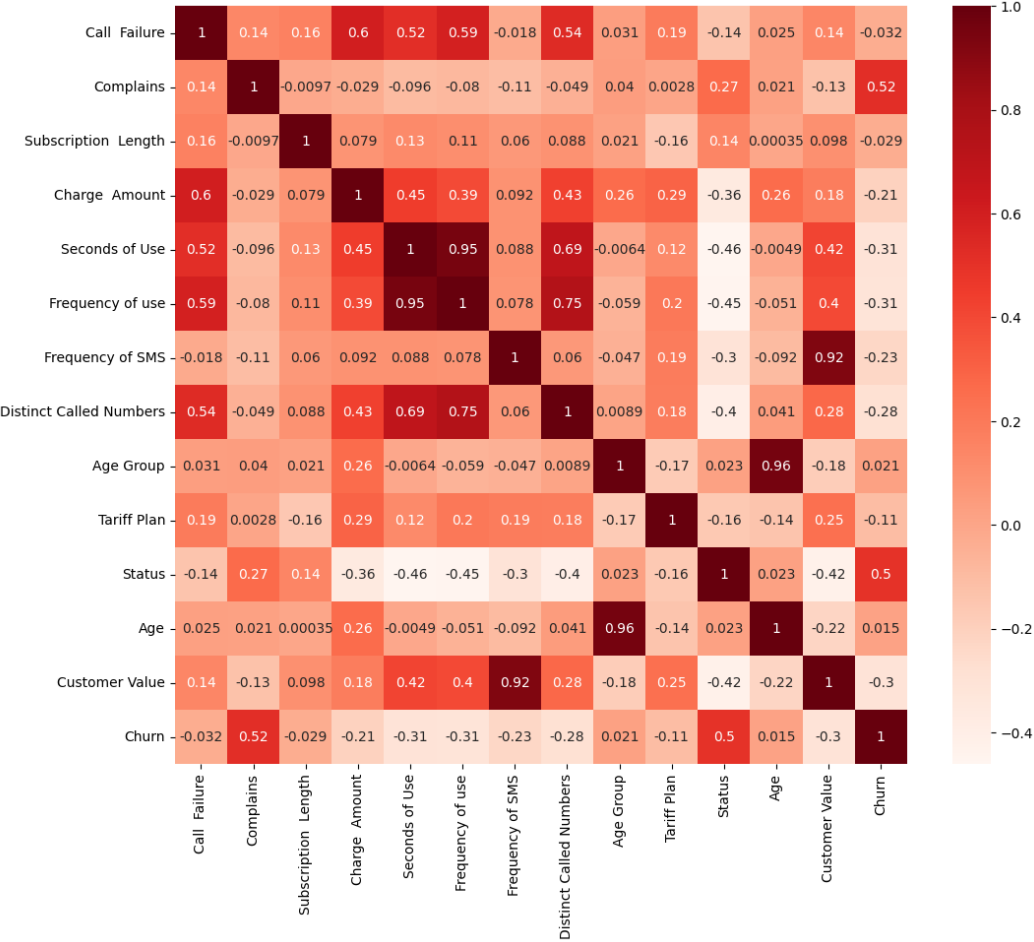
ROC AUC: 0.7655255112480591

```
# Cross-validation
for name, model in models.items():
    scores = cross_val_score(model, X_train, y_train, cv=5, scoring='roc_auc')
    print(f"Model: {name}, ROC AUC: {scores.mean()}")

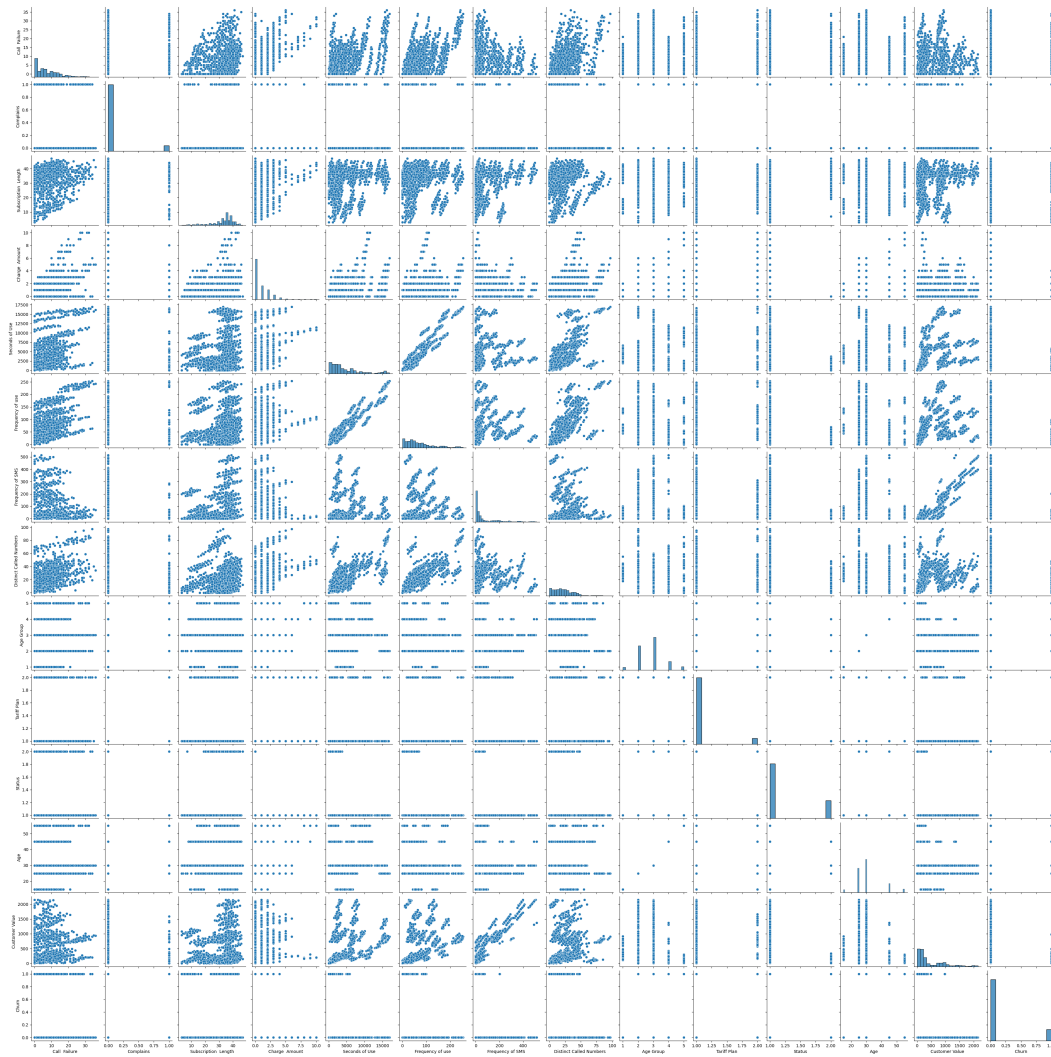
# Confusion Matrix
for name, model in models.items():
    y_pred = model.predict(X_test)
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt="d")
    plt.title(f"Confusion matrix for {name}")
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.show()
```



```
plt.figure(figsize=(12,10))
corr = train_data.corr()
sns.heatmap(corr, annot=True, cmap=plt.cm.Reds)
plt.show()
```




```
sns.pairplot(train_data)
plt.show()
```



```

model = RandomForestClassifier()
model.fit(X_train, y_train)
importance = model.feature_importances_
plt.bar([x for x in range(len(importance))], importance)
plt.show()

```



```
print(train_data.describe())
print(test_data.describe())
```

	Distinct Called Numbers	Age Group	Tariff Plan	Status \
count	1575.000000	1575.000000	1575.000000	1575.000000
mean	23.537143	2.806349	1.079365	1.245079
std	17.718890	0.873853	0.270394	0.430271
min	0.000000	1.000000	1.000000	1.000000
25%	10.000000	2.000000	1.000000	1.000000
50%	21.000000	3.000000	1.000000	1.000000
75%	33.000000	3.000000	1.000000	1.000000
max	97.000000	5.000000	2.000000	2.000000

	Age	Customer Value	Churn
count	1575.000000	1575.000000	1575.000000
mean	30.780952	480.649708	0.154921
std	8.586189	515.837852	0.361944
min	15.000000	0.000000	0.000000
25%	25.000000	114.412500	0.000000
50%	30.000000	237.915000	0.000000
75%	30.000000	807.727500	0.000000
max	55.000000	2149.280000	1.000000



	Call Failure	Complains	Subscription Length	Charge Amount \
count	1575.000000	1575.000000	1575.000000	1575.000000
mean	7.589841	0.075556	32.845079	0.944127
std	7.213174	0.264370	8.494867	1.557268
min	0.000000	0.000000	3.000000	0.000000
25%	1.000000	0.000000	30.000000	0.000000
50%	6.000000	0.000000	35.000000	0.000000
75%	11.500000	0.000000	38.000000	1.000000
max	36.000000	1.000000	46.000000	10.000000

	Seconds of Use	Frequency of use	Frequency of SMS \
count	1575.000000	1575.000000	1575.000000
mean	4378.377778	68.247619	72.452063
std	4078.714808	55.752091	113.154851
min	0.000000	0.000000	0.000000
25%	1366.500000	27.000000	7.000000
50%	2970.000000	54.000000	20.000000
75%	6480.000000	94.000000	81.500000
max	17090.000000	255.000000	522.000000




	Distinct Called Numbers	Age Group	Tariff Plan	Status \
count	1575.000000	1575.000000	1575.000000	1575.000000
mean	23.482540	2.845714	1.076190	1.251429
std	16.706322	0.910726	0.265387	0.433972
min	0.000000	1.000000	1.000000	1.000000
25%	10.000000	2.000000	1.000000	1.000000
50%	21.000000	3.000000	1.000000	1.000000
75%	34.000000	3.000000	1.000000	2.000000
max	88.000000	5.000000	2.000000	2.000000

	Age	Customer Value	Churn
count	1575.000000	1575.000000	1575.000000
mean	31.215873	461.296124	0.159365
std	9.066905	518.173380	0.366132
min	15.000000	0.000000	0.000000
25%	25.000000	111.760000	0.000000
50%	30.000000	221.715000	0.000000
75%	30.000000	770.715000	0.000000
max	55.000000	2165.280000	1.000000

```
plt.figure(figsize=(8,6))
sns.countplot(x='Churn', data=train_data)
plt.show()
plt.figure(figsize=(8,6))
sns.countplot(x='Churn', data=train_data)
plt.show()
```



```
for column in train_data.columns:
    if column != 'Churn':
        plt.figure(figsize=(8,6))
        sns.histplot(data=train_data, x=column, hue='Churn', element='step', stat='density', common_norm=False)
        plt.show()
```


```
from sklearn.model_selection import cross_val_score

for name, model in models.items():
    cv_scores = cross_val_score(model, X_train, y_train, cv=5)
    print(f"{name} Cross Validation Score: {cv_scores.mean()}")

    Logistic Regression Cross Validation Score: 0.8920634920634921
    Random Forest Cross Validation Score: 0.947936507936508
    Support Vector Machine Cross Validation Score: 0.9117460317460317
```

```
# Import necessary libraries
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import pandas as pd
```

```
# Prepare the training data
X_train = train_data.drop('Churn', axis=1)
y_train = train_data['Churn']
```

```
# Prepare the test data
X_test = test_data.drop('Churn', axis=1)
y_test = test_data['Churn']
```

```
# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
# Train a logistic regression model
model = LogisticRegression()
model.fit(X_train, y_train)
```

```
# Make predictions on the test data
y_pred = model.predict(X_test)
```

```
# Print a classification report
print(classification_report(y_test, y_pred))
```

```
# Print a confusion matrix
print(confusion_matrix(y_test, y_pred))
```

```

              precision    recall  f1-score   support

     0       0.90      0.97      0.94      1324
     1       0.76      0.43      0.55       251

 accuracy      0.83
 macro avg     0.83
 weighted avg   0.83

```

```
# Import necessary libraries
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc
```

```
# Prepare the training data
X_train = train_data.drop('Churn', axis=1)
y_train = train_data['Churn']
```

```
# Prepare the test data
X_test = test_data.drop('Churn', axis=1)
y_test = test_data['Churn']
```

```
# Standardize the features
scaler = StandardScaler()
```